


All Pair Shortest Path Algorithm

Floyd-Warshall

Adapted from
Stanford CS161 course

Recall

- A weighted directed graph:



- Weights on edges represent **costs**.
- The **cost of a path** is the sum of the weights along that path.
- A **shortest path** from s to t is a directed path from s to t with the smallest cost.
- The **single-source shortest path problem** is to find the shortest path from s to v for all v in the graph.

Last time

- Dijkstra's algorithm!
 - Solves the single-source shortest path problem in weighted graphs.
- Bellman-Ford algorithm!
 - **ALSO** solves the single-source shortest path problem in weighted graphs.

Recap: shortest paths

- **BFS:**

- (+) $O(n+m)$
- (-) only unweighted graphs

- **Dijkstra's algorithm:**

- (+) weighted graphs
- (+) $O(n \log(n) + m)$ if you implement it with a Fibonacci heap
- (-) no negative edge weights
- (-) very “centralized” (need to keep track of all the vertices to know which to update).

- **Bellman-Ford algorithm:**


- (+) weighted graphs, even with negative weights
- (+) can be done in a distributed fashion, every vertex using only information from its neighbors.
- (-) $O(nm)$

Bellman-Ford vs. Dijkstra

- Dijkstra:
 - Find the u with the smallest $d[u]$
 - Update u 's neighbors: $d[v] = \min(d[v], d[u] + w(u,v))$
- Bellman-Ford:
 - Don't bother finding the u with the smallest $d[u]$
 - Everyone updates!
 - Slower, but more flexible:
 - Can handle negative edge weights (as long as there aren't negative cycles)
 - Can do updates in a decentralized way.

Aside: Negative Cycles

- A **negative cycle** is a cycle whose edge weights sum to a negative number.
- Shortest paths aren't defined when there are negative cycles!



The shortest path from A to B has cost...negative infinity?

Bellman-Ford vs. Dijkstra

- Dijkstra:
 - Find the u with the smallest $d[u]$
 - Update u 's neighbors: $d[v] = \min(d[v], d[u] + w(u,v))$
- Bellman-Ford:
 - Don't bother finding the u with the smallest $d[u]$
 - Everyone updates!
 - Slower, but more flexible:
 - Can handle negative edge weights (as long as there aren't negative cycles)
 - Can do updates in a decentralized way.

$G = (V, E)$ is a graph with n vertices and m edges.

Bellman-Ford* algorithm

Bellman-Ford*(G,s):

- Initialize arrays $d^{(0)}, \dots, d^{(n-1)}$ of length n
- $d^{(0)}[v] = \infty$ for all v in V
- $d^{(0)}[s] = 0$
- **For** $i=0, \dots, n-2$:
 - **For** v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], \min_{u \text{ in } v.\text{inNbrs}}\{d^{(i)}[u] + w(u, v)\})$
- Now, $\text{dist}(s, v) = d^{(n-1)}[v]$ for all v in V .
 - (Assuming no negative cycles)


Here, Dijkstra picked a special vertex u and updated u 's neighbors – Bellman-Ford will update all the vertices.



*Slightly different than some versions of Bellman-Ford...but
this way is pedagogically convenient for today's lecture.⁸

Note on implementation


- Don't actually keep all n arrays around.
- Just keep two at a time: "last round" and "this round"



Interpretation of $d^{(i)}$

For all vertices v , $d^{(i)}[v]$ is equal to the cost of the shortest path between s and v with at most i edges.


	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$	0	1	2	6	23



Why does Bellman-Ford work?

- Inductive hypothesis:
 - For all v , $d^{(i)}[v]$ is equal to the cost of the shortest path between s and v **with at most i edges.**
- Conclusion:
 - For all v , $d^{(n-1)}[v]$ is equal to the cost of the shortest path between s and v **with at most $n-1$ edges.**

Do the base case and
inductive step!




Aside: simple paths

Assume there is no negative cycle.

“Simple” means that the path has no cycles in it.




- Then there is a shortest path from s to t , and moreover there is a **simple** shortest path.



This cycle isn't helping.
Just get rid of it.

- A **simple path** in a graph with n vertices has at most $n-1$ edges in it.

Can't add another edge without making a cycle!



- So there is a shortest path with at most $n-1$ edges

Why does Bellman-Ford work?

- Inductive hypothesis:
 - For all v , $d^{(i)}[v]$ is equal to the cost of the shortest path between s and v with at most i edges.
- Conclusion:
 - For all v , $d^{(n-1)}[v]$ is equal to the cost of the shortest path between s and v with at most $n-1$ edges.
 - If there are no negative cycles, $d^{(n-1)}[v]$ is equal to the cost of the shortest path.

Notice that negative edge weights are fine.
Just not negative cycles.

Bellman-Ford take-aways

- Running time is $O(mn)$
 - For each of n rounds, update m edges.
- Works fine with negative edges.
- Does not work with negative cycles.
 - No algorithm can – shortest paths aren't defined if there are negative cycles.
- B-F can detect negative cycles!
 - Fun exercise! (Hint, what happens if you run Bellman-Ford for longer than $n-1$ iterations?)

- What class of algorithm does Dijkstra fall in?

Greedy Algorithm

- What about Bellman-Ford?

Bellman-Ford is an example of...

Dynamic Programming!

What about Folyd-Warshall??


Also, Dynamic Programming

Candidate algorithm

- **def** Fibonacci (n) :
 - **if** n == 0, **return** 0
 - **if** n == 1, **return** 1
 - **return** Fibonacci (n-1) + Fibonacci (n-2)


Running time?

- $T(n) = T(n-1) + T(n-2) + O(1)$
- $T(n) \geq T(n-1) + T(n-2)$ for $n \geq 2$
- So $T(n)$ grows *at least* as fast as the Fibonacci numbers themselves...
- This is **EXPONENTIALLY QUICKLY!**




What's going on?
Consider $\text{Fib}(8)$

That's a lot of
repeated
computation!




Maybe this would be better:



```
def fasterFibonacci(n):
    • F = [0, 1, None, None, ..., None]
        • \\\ F has length n + 1
    • for i = 2, ..., n:
        • F[i] = F[i-1] + F[i-2]
    • return F[n]
```

Much better running time!



This was an example of...

*Dynamic
Programming!*

What is *dynamic programming*?

- It is an algorithm design paradigm
 - like divide-and-conquer is an algorithm design paradigm.
- Usually it is for solving **optimization problems**
 - eg, *shortest* path
 - (Fibonacci numbers aren't an optimization problem, but they are a good example of DP anyway...)

Elements of dynamic programming

1. Optimal sub-structure:

- Big problems break up into sub-problems.
 - Fibonacci: $F(i)$ for $i \leq n$
 - Bellman-Ford: Shortest paths with at most i edges for $i < n$
- The optimal solution to a problem can be expressed in terms of optimal solutions to smaller sub-problems.
 - Fibonacci:

$$F(i+1) = F(i) + F(i-1)$$

- Bellman-Ford:

$$d^{(i+1)}[v] \leftarrow \min\{ d^{(i)}[v], \min_u \{d^{(i)}[u] + \text{weight}(u,v)\} \}$$

Shortest path with at
most i edges from s to v

Shortest path with at most
 i edges from s to u .

*The word “optimal” makes sense in the context of optimization problems like shortest path, and is why this is called “Optimal Sub-structure.”

Elements of dynamic programming

2. Overlapping sub-problems:

- The sub-problems overlap.
 - Fibonacci:
 - Both $F[i+1]$ and $F[i+2]$ directly use $F[i]$.
 - And lots of different $F[i+x]$ indirectly use $F[i]$.
 - Bellman-Ford:
 - Many different entries of $d^{(i+1)}$ will directly use $d^{(i)}[v]$.
 - And lots of different entries of $d^{(i+x)}$ will indirectly use $d^{(i)}[v]$.
- This means that we can save time by solving a sub-problem just once and storing the answer.

Elements of dynamic programming


- Optimal substructure.
 - Optimal solutions to sub-problems can be used to find the optimal solution of the original problem.
- Overlapping subproblems.
 - The subproblems show up again and again
- Using these properties, we can design a ***dynamic programming*** algorithm:
 - Keep a table of solutions to the smaller problems.
 - Use the solutions in the table to solve bigger problems.
 - At the end we can use information we collected along the way to find the solution to the whole thing.

Two ways to think about and/or implement DP algorithms

- Top down
- Bottom up

Bottom up approach


what we just saw.



- For Fibonacci:
- Solve the small problems first
 - fill in $F[0], F[1]$
- Then bigger problems
 - fill in $F[2]$
- ...
- Then bigger problems
 - fill in $F[n-1]$
- Then finally solve the real problem.
 - fill in $F[n]$


Bottom up approach

what we just saw.



- For Bellman-Ford:
- Solve the small problems first
 - fill in $d^{(0)}$
- Then bigger problems
 - fill in $d^{(1)}$
- ...
- Then bigger problems
 - fill in $d^{(n-2)}$
- Then finally solve the real problem.
 - fill in $d^{(n-1)}$

Top down approach




- Think of it like a recursive algorithm.
- To solve the big problem:
 - Recurse to solve smaller problems
 - Those recurse to solve smaller problems
 - etc..
- The difference from divide and conquer:
 - Keep track of what small problems you've already solved to prevent re-solving the same problem twice.
 - Aka, “**memo-ization**”



Example of top-down Fibonacci

- define a global list $F = [0, 1, \text{None}, \text{None}, \dots, \text{None}]$
- **def** Fibonacci(n):
 - **if** $F[n] \neq \text{None}$:
 - **return** $F[n]$
 - **else**:
 - $F[n] = \text{Fibonacci}(n-1) + \text{Fibonacci}(n-2)$
 - **return** $F[n]$

Memo-ization:
Keeps track (in F) of
the stuff you've
already done.



What have we learned?

- ***Dynamic programming:***


- Paradigm in algorithm design.
- Uses **optimal substructure**
- Uses **overlapping subproblems**
- Can be implemented **bottom-up** or **top-down**.
- It's a fancy name for a pretty common-sense idea:



Floyd-Warshall Algorithm

- This is an algorithm for **All-Pairs Shortest Paths (APSP)**
 - That is, I want to know the shortest path from u to v for **ALL pairs** u,v of vertices in the graph.
 - Not just from a special single source s .

Destination		s	u	v	t
Source	s	0	2	4	2
u	1	0	2	0	
v	∞	∞	0	-2	
t	∞	∞	∞	0	



Floyd-Warshall Algorithm

Another example of DP

- This is an algorithm for **All-Pairs Shortest Paths (APSP)**
 - That is, I want to know the shortest path from u to v for **ALL pairs** u,v of vertices in the graph.
 - Not just from a special single source s .
- Naïve solution (if we want to handle negative edge weights):
 - For all s in G :
 - Run Bellman-Ford on G starting at s .
 - Time $O(n \cdot nm) = O(n^2m)$,
 - may be as bad as n^4 if $m=n^2$

Can we do better?


Recipe for applying Dynamic Programming

- **Step 1:** Identify optimal substructure.
 - What are our subproblems?
- **Step 2:** Find a recursive formulation for the subproblems
 - How can we solve larger problems using smaller ones?
- **Step 3:** Use dynamic programming to find the thing you want.
 - Fill in a table, starting with the smallest sub-problems and building up.



Label the vertices 1,2,...,n

Optimal substructure




Optimal substructure

Label the vertices $1, 2, \dots, n$
(We omit some edges in the picture below – meant to be a cartoon, not an example).

Sub-problem($k-1$):

For all pairs, u, v , find the cost of the shortest path from u to v , so that all the internal vertices on that path are in $\{1, \dots, k-1\}$.

Let $D^{(k-1)}[u, v]$ be the solution to Sub-problem($k-1$).




Recipe for applying Dynamic Programming

- **Step 1:** Identify optimal substructure.
 - What are our subproblems?
 - **Step 2:** Find a recursive formulation for the subproblems
 - How can we solve larger problems using smaller ones?
 - **Step 3:** Use dynamic programming to find the thing you want.
 - Fill in a table, starting with the smallest sub-problems and building up.
- 

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?


$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.


Case 1: we don't need vertex k .



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.

Case 2: we need vertex k .




Case 2 continued

- Suppose there are no negative cycles. WLOG the shortest path from u to v through $\{1, \dots, k\}$ is simple.

- The shortest path from u to v looks like this:

- This path is the shortest path from u to k through $\{1, \dots, k-1\}$.
 - sub-paths of shortest paths are shortest paths
- Similarly for this path.


Case 2: we need vertex k .



$$D^{(k)}[u, v] = D^{(k-1)}[u, k] + D^{(k-1)}[k, v]$$


How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

Case 1: we don't need vertex k .



$$D^{(k)}[u,v] = D^{(k-1)}[u,v]$$

Case 2: we need vertex k .



$$D^{(k)}[u,v] = D^{(k-1)}[u,k] + D^{(k-1)}[k,v]$$

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

Case 1: Cost of
shortest path
through $\{1,\dots,k-1\}$

Case 2: Cost of shortest path
from **u** to **k** and then from **k** to **v**
through $\{1,\dots,k-1\}$

- Optimal substructure:
 - We can solve the big problem using solutions to smaller problems.
- Overlapping sub-problems:
 - $D^{(k-1)}[k,v]$ can be used to help compute $D^{(k)}[u,v]$ for lots of different u 's.

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

Case 1: Cost of
shortest path
through $\{1,\dots,k-1\}$

Case 2: Cost of shortest path
from **u** to **k** and then from **k** to **v**
through $\{1,\dots,k-1\}$

- Using our ***Dynamic programming*** paradigm, this gives us an algorithm!



Recipe for applying Dynamic Programming

- **Step 1:** Identify optimal substructure.
 - What are our subproblems?
 - **Step 2:** Find a recursive formulation for the subproblems
 - How can we solve larger problems using smaller ones?
 - **Step 3:** Use dynamic programming to find the thing you want.
 - Fill in a table, starting with the smallest sub-problems and building up.
- 

Floyd-Warshall algorithm

- Initialize n -by- n arrays $D^{(k)}$ for $k = 0, \dots, n$

- $D^{(0)}[u,v] = \infty$ for all pairs (u,v)

- $D^{(0)}[u,u] = 0$ for all u

- $D^{(0)}[u,v] = \text{weight}(u,v)$ for all (u,v) in E .

The base case checks out: the only path through zero other vertices are edges directly from u to v .

- For $k = 1, \dots, n$:

- For pairs u,v in V^2 :

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

- Return $D^{(n)}$

This is a bottom-up *Dynamic programming* algorithm.

Our earlier logic shows

- Theorem:
 - If there are no negative cycles in a weighted directed graph G , then the Floyd-Warshall algorithm, running on G , returns a matrix $D^{(n)}$ so that:

$D^{(n)}[u,v] = \text{distance between } u \text{ and } v \text{ in } G.$

- Running time: $O(n^3)$
 - Better than running Bellman-Ford n times!
- Storage:
 - Need to store **two** n -by- n arrays, and the original graph.

Work out the details of a proof!



As with Bellman-Ford, we don't really need to store all n of the $D^{(k)}$.

What if there *are* negative cycles?

- Just like Bellman-Ford, Floyd-Warshall can detect negative cycles:
 - “Negative cycle” means that there’s some v so that there is a path from v to v that has cost < 0 .
 - Aka, $D^{(n)}[v,v] < 0$.
- Algorithm:
 - Run Floyd-Warshall as before.
 - If there is some v so that $D^{(n)}[v,v] < 0$:
 - **return** negative cycle.

Can we do better than $O(n^3)$?

Nothing on this slide is required knowledge for this class

- There is an algorithm that runs in time $O(n^3/\log^{100}(n))$.
 - [Williams, “Faster APSP via Circuit Complexity”, STOC 2014]
- If you can come up with an algorithm for All-Pairs-Shortest-Path that runs in time $O(n^{2.99})$, that would be a really big deal.
 - Let me know if you can!
 - See [Abboud, Vassilevska-Williams, “Popular conjectures imply strong lower bounds for dynamic problems”, FOCS 2014] for some evidence that this is a very difficult problem!

Recipe for applying Dynamic Programming

- **Step 1:** Identify optimal substructure.
- **Step 2:** Find a recursive formulation for the thing you want.
 - E.g, length of shortest paths
- **Step 3:** Use dynamic programming to find the thing you want.
 - Fill in a table, starting with the smallest sub-problems and building up.
- **(Steps 4 and 5 coming next lecture...)**