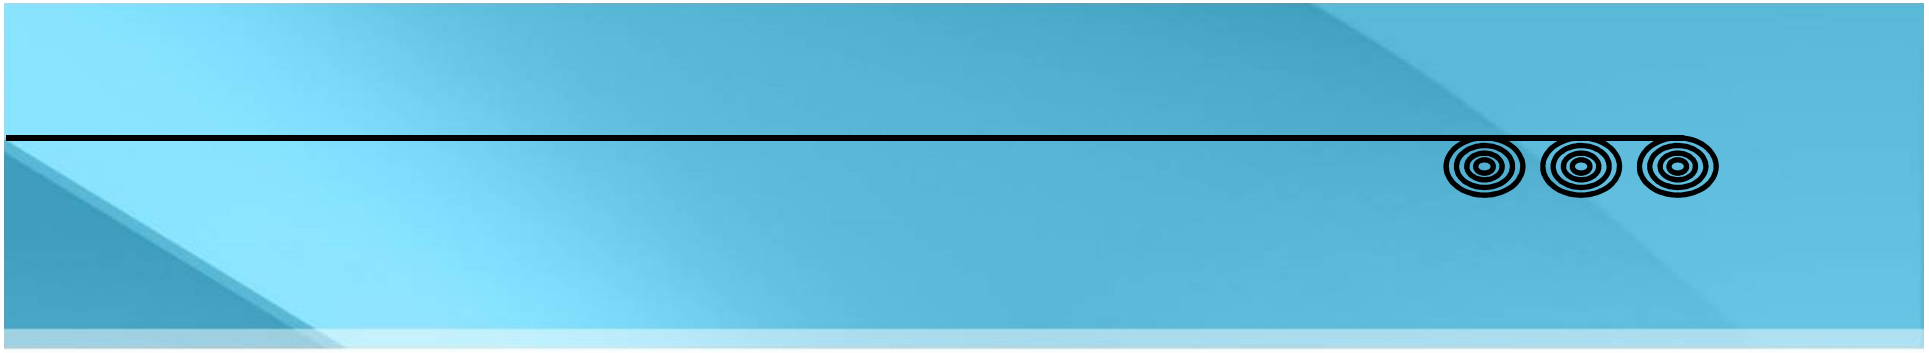

CSE 215: Database

Department of CSE
Bangladesh University of Engineering & Technology

What is it?

- The ER data model employs three basic concepts:
  - entity sets
  - relationship sets
  - attributes
- The ER model also has an associated diagrammatic representation, which can express the overall logical structure of a database graphically
  - ER diagram

Entity Sets

- An **entity** is an object that exists and is distinguishable from other objects
  - Example: specific person, company, plant
- An **entity set** is a set of entities of the same type that share the same properties
  - Example: set of all persons, companies, trees,
- An entity is represented by a set of attributes; i.e., descriptive properties possessed by all members of an entity set
  - Example:

    instructor = (ID, name, salary)
    course= (course_id, title, credits)
- A subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set

Entity Sets – ER Diagram

- Entity sets can be represented graphically as follows:
  - Rectangles represent entity sets.
  - Attributes listed inside entity rectangle
  - Underline indicates primary key attributes

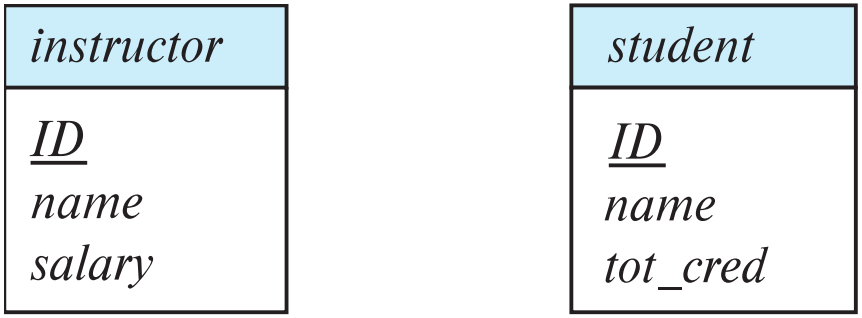

Relationship Sets

- A **relationship** is an association among several entities

  Example:

  44553 (Peltier) *advisor* 22222 (Einstein)

  student entity relationship set *instructor* entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \ldots e_n) \mid e_1 \in E_1, e_2 \in E_2, \ldots, e_n \in E_n\}$$

  where $(e_1, e_2, \ldots, e_n)$ is a relationship

  – Example:

$$(44553,22222) \in advisor$$

Relationship Sets

- A **relationship** is an association among several entities

  Example:

  44553 (Peltier) *advisor* 22222 (Einstein)

  student entity relationship set *instructor* entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \ldots e_n) \mid e_1 \in E_1, e_2 \in E_2, \ldots, e_n \in E_n\}$$

  where $(e_1, e_2, \ldots, e_n)$ is a relationship

  – Example:

$$(44553,22222) \in advisor$$

Relationship Sets

- Example: we define the relationship set advisor to denote the associations between students and the instructors who act as their advisors
- Pictorially, we draw a line between related entities

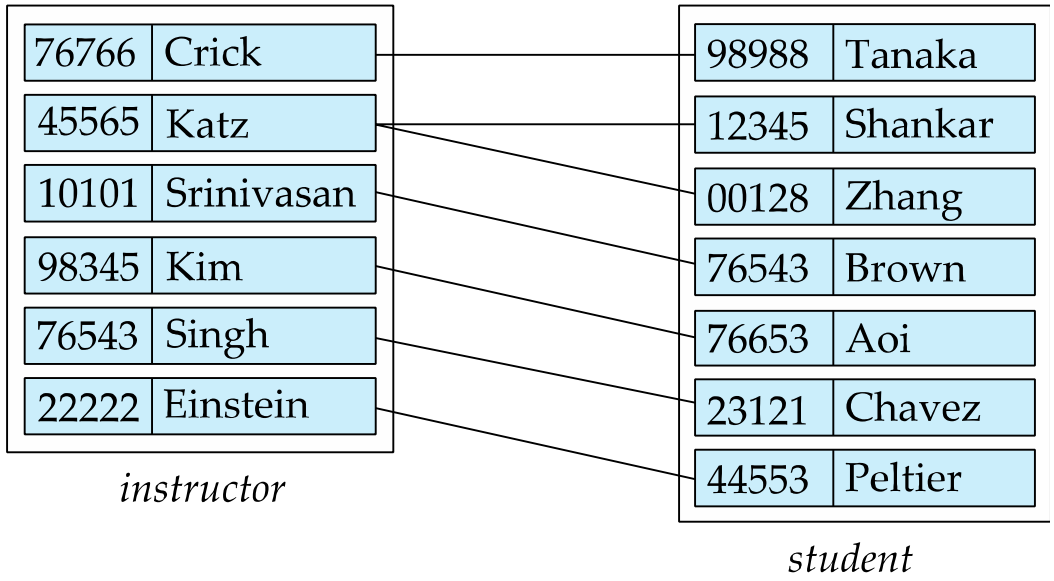

Relationship Sets-ER Diagram

- Diamonds represent relationship sets

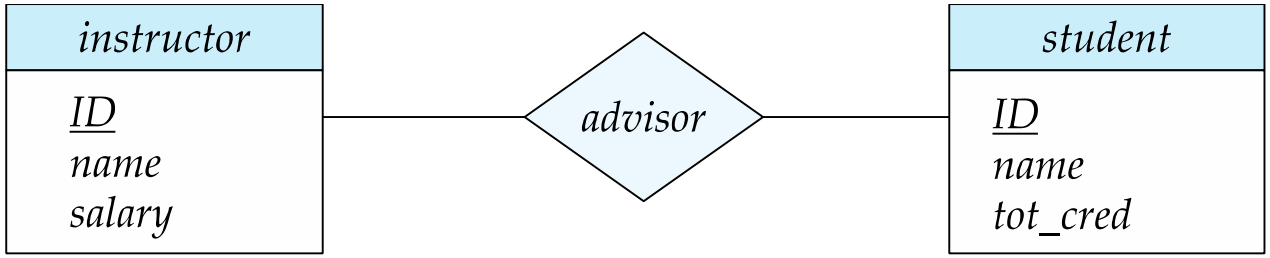

Relationship Sets

- An attribute can also be associated with a relationship set.
- For instance, the advisor relationship set between entity sets instructor and student may have the attribute date which tracks when the student started being associated with the advisor

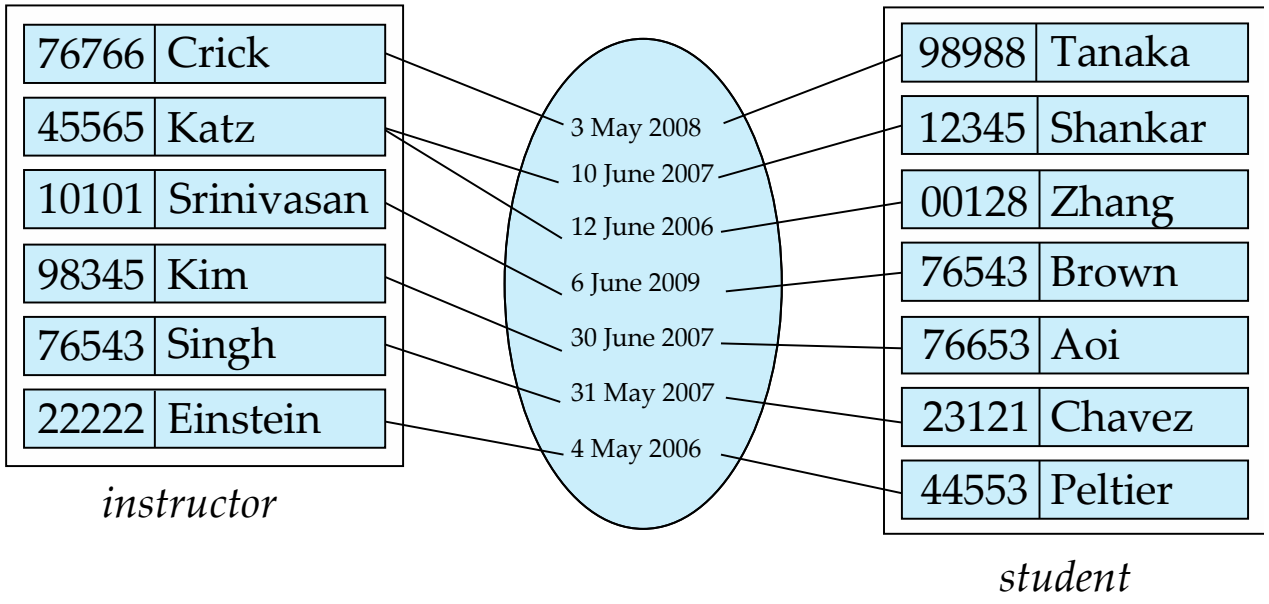

Relationship Sets-ER Diagram

- Diamonds represent relationship sets

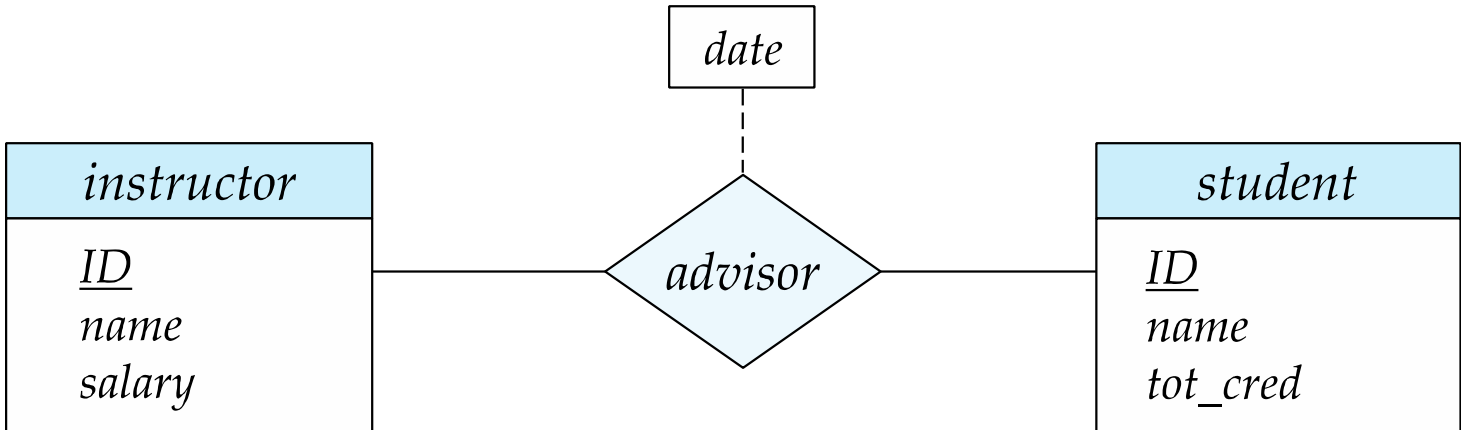

Roles

- Entity sets of a relationship need not be distinct
  - Each occurrence of an entity set plays a "**role**" in the relationship
- The labels "course_id" and "prereq_id" are called roles.

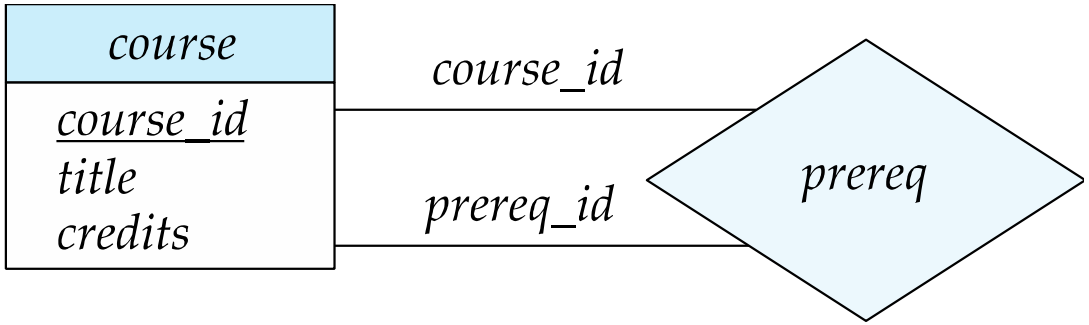

Relationship Degree

- Binary relationship
  - Involve two entity sets (or degree two)
- Relationships between more than two entity sets are rare
  - Example: students work on research projects under the guidance of an instructor.
  - relationship proj_guide is a ternary relationship between instructor, student, and project

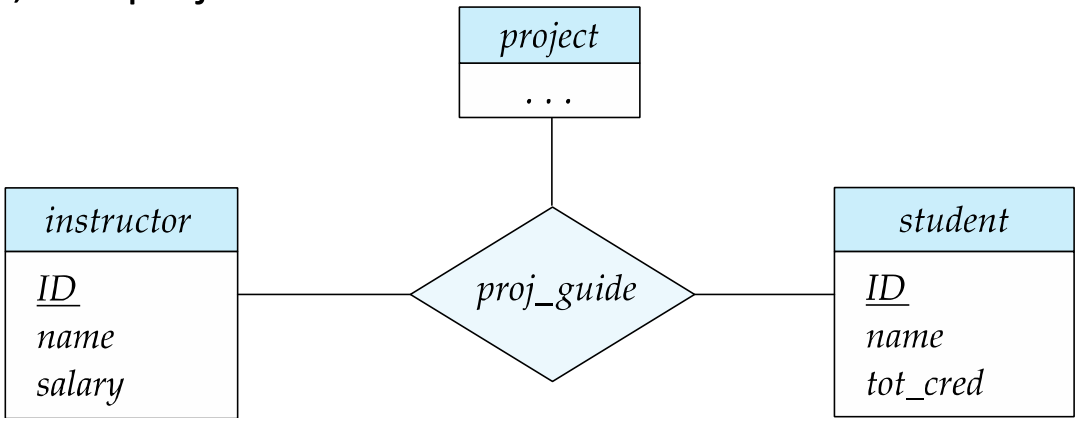

Attributes

- Attribute types:
  - **Simple** and **composite** attributes
    - Composite attributes allow us to divided attributes into subparts
  - **Single-valued** and **multivalued** attributes
    - Example: multivalued attribute: *phone_numbers*
  - **Derived** attributes
    - Can be computed from other attributes
    - Example: age, given date_of_birth
- **Domain** – the set of permitted values for each attribute

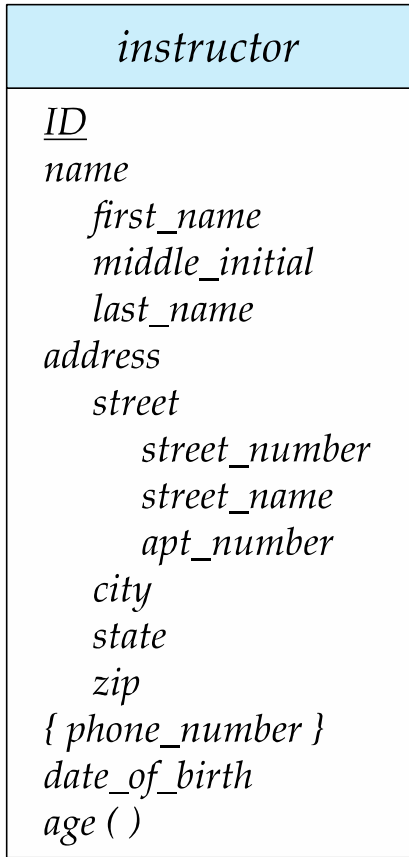

Mapping Cardinalities

- Express the number of entities to which another entity can be associated via a relationship set
- **One to one**
  - An entity in A is associated with at most one entity in B, An entity in B is associated with at most one entity in A
- **One to many**
  - An entity in A is associated with any number (0 or more) of entities in B, An entity in B is associated with at most one entity in A
- **Many to one**
  - An entity in A is associated with at most one entity in B, An entity in B is associated with any number (0 or more) of entities in A
- **Many to many**
  - An entity in A is associated with any number (0 or more) of entities in B, An entity in B is associated with any number (0 or more) of entities in A

Mapping Cardinalities

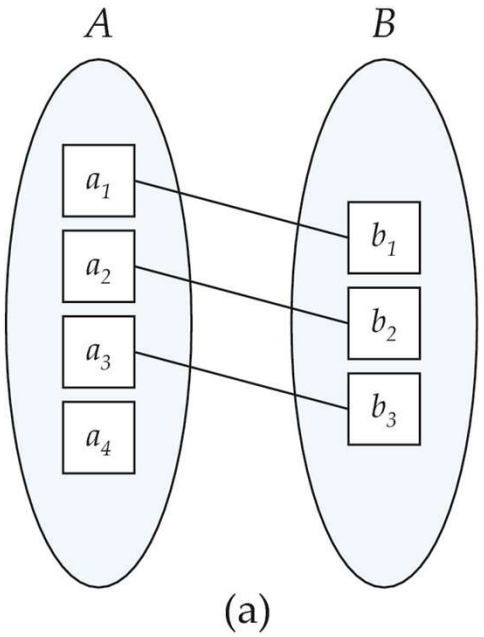

One to one

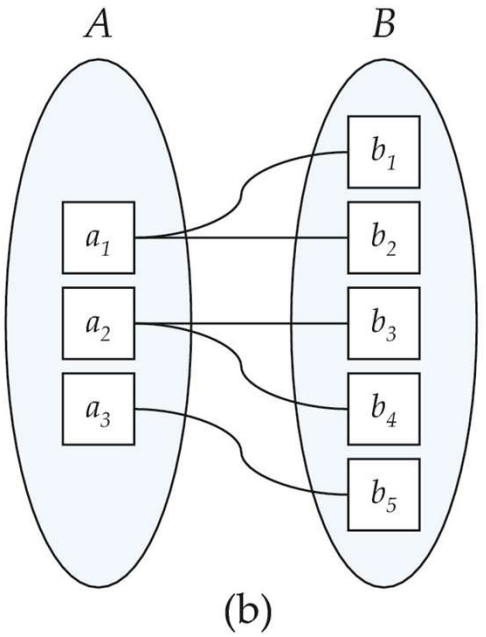

One to many

Mapping Cardinalities

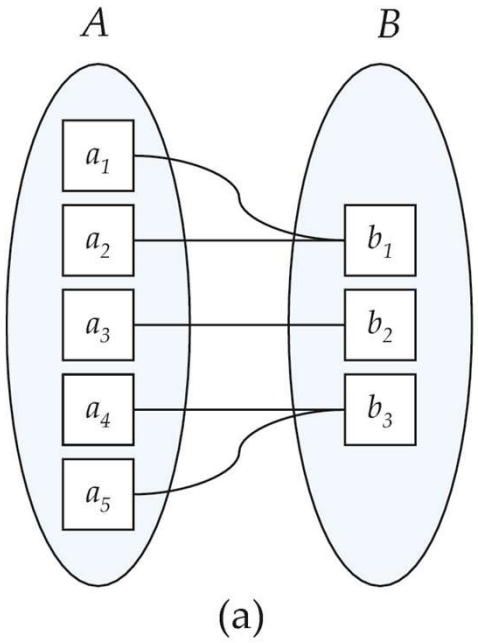

Many to one

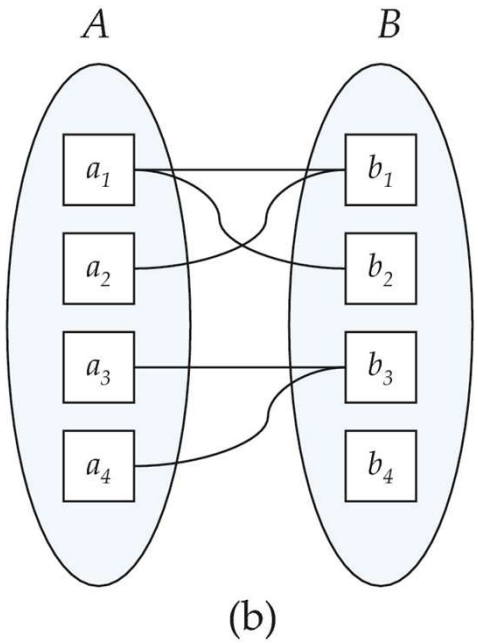

Many to many

Cardinality Constraints-ER Diagram

- We express cardinality constraints by drawing either a directed line (□), signifying "one," or an undirected line (—), signifying "many," between the relationship set and the entity set

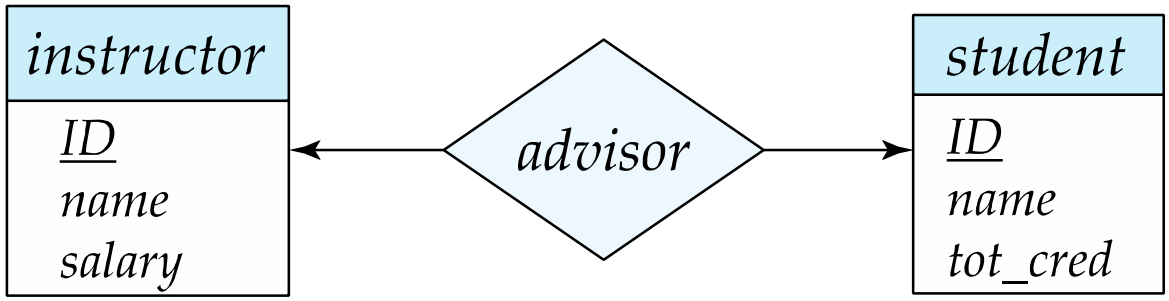

One to one

Cardinality Constraints-ER Diagram

- We express cardinality constraints by drawing either a directed line (□), signifying "one," or an undirected line (—), signifying "many," between the relationship set and the entity set

One to many

Cardinality Constraints-ER Diagram

- We express cardinality constraints by drawing either a directed line (□), signifying "one," or an undirected line (—), signifying "many," between the relationship set and the entity set

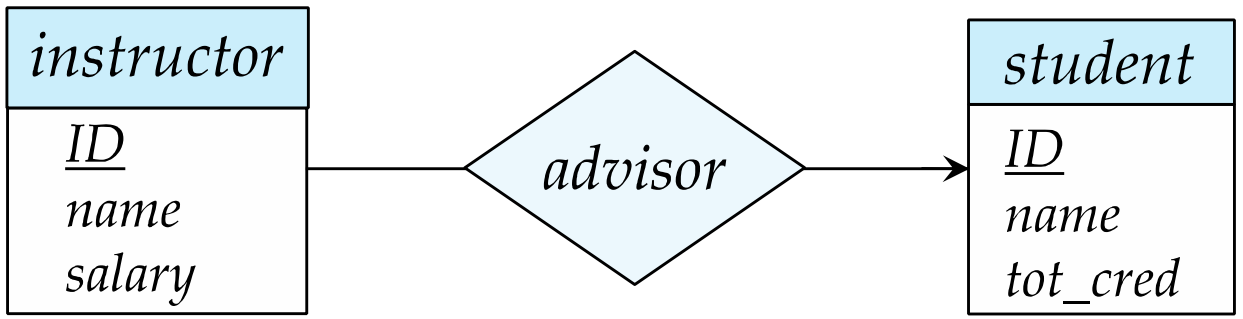

Many to one

Cardinality Constraints-ER Diagram

- We express cardinality constraints by drawing either a directed line (□), signifying "one," or an undirected line (—), signifying "many," between the relationship set and the entity set

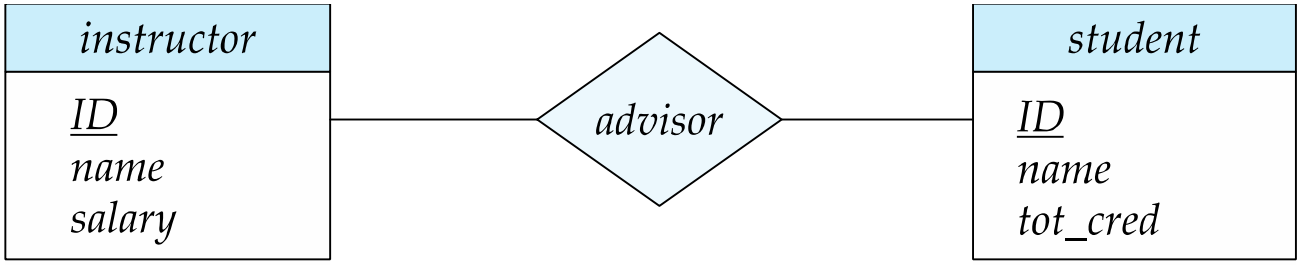

Many to many

Cardinality Constraints-ER Diagram

- **Total participation** (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
  - Example: participation of student in advisor relation is total, every student must have an associated instructor

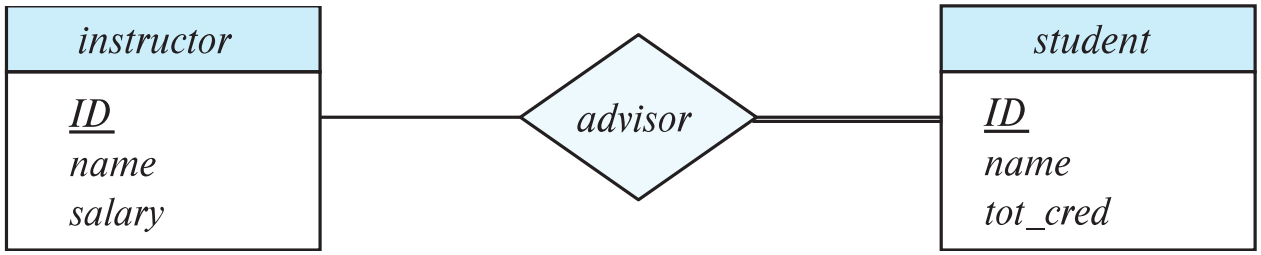

- **Partial participation**: some entities may not participate in any relationship in the relationship set
  - Example: participation of instructor in advisor is partial

Cardinality Constraints-ER Diagram

- A line may have an associated minimum and maximum cardinality, shown in the form l..h, where l is the minimum and h the maximum cardinality
  - A minimum value of 1 indicates total participation
  - A maximum value of 1 indicates that the entity participates in at most one relationship
  - A maximum value of * indicates no limit

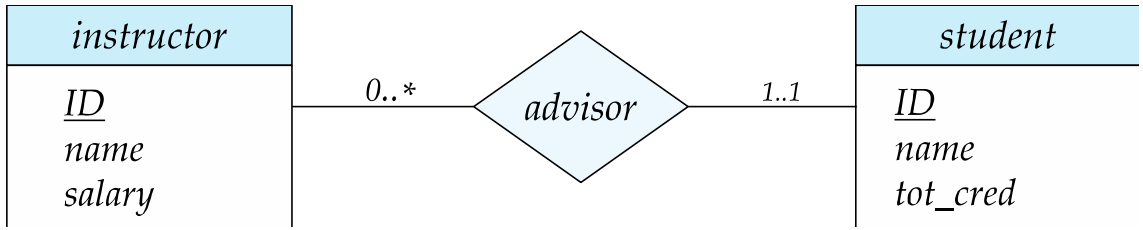

Instructor can advise 0 or more students
A student must have 1 advisor; cannot have multiple advisors

Cardinality Constraints-ER Diagram

- We allow at most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint
  - For example, an arrow from proj_guide to instructor indicates each student has at most one guide for a project
- If there is more than one arrow, there are two ways of defining the meaning.
  - For example, a ternary relationship R between A, B and C with arrows to B and C could mean
    1. Each A entity is associated with a unique entity from B and C
    2. Each pair of entities from (A, B) is associated with a unique C entity, and each pair (A, C) is associated with a unique B

Primary Key for Entity Sets

- By definition, individual entities are distinct.
- From database perspective, the differences among them must be expressed in terms of their attributes.

- The values of the attribute values of an entity must be such that they can uniquely identify the entity.
  - No two entities in an entity set are allowed to have exactly the same value for all attributes.

- A **key** for an entity is a set of attributes that suffice to distinguish entities from each other

Keys

- Let K is a subset of attributes
- K is a superkey if values for K are sufficient to identify a unique tuple of each possible relation r(R)
  - Example: {ID} and {ID,name} are both superkeys of instructor
- Superkey K is a candidate key if K is minimal
  - Example: {ID} is a candidate key for Instructor
- One of the candidate keys is selected to be the primary key

Primary Key for Relationship Sets

- To distinguish among the various relationships of a relationship set we use the individual primary keys of the entities in the relationship set

- Let R be a relationship set involving entity sets E1, E2, .. En
- The primary key for R is consists of the union of the primary keys of entity sets E1, E2, ..En
  - If the relationship set R has attributes a1, a2, .., am associated with it, then the primary key of R and attributes a1, a2, .., am describe the relationship in set R
- Example: relationship set “advisor”
  - The primary key consists of instructor.ID and student.ID

Primary Key for Relationship Sets

- The choice of the primary key for a relationship set depends on the mapping cardinality of the relationship set
- Many-to-Many relationships
  - The preceding union of the primary keys is a minimal superkey and is chosen as the primary key
- One-to-Many relationships
  - The primary key of the "Many" side is a minimal superkey and is used as the primary key
- Many-to-one relationships
  - The primary key of the "Many" side is a minimal superkey and is used as the primary key
- One-to-one relationships
  - The primary key of either one of the participating entity sets forms a minimal superkey, and either one can be chosen as the primary key

Weak/Strong Entity Sets

- A **weak entity set** is one whose existence is dependent on another entity, called its **identifying entity**
- Instead of associating a primary key with a weak entity, we use the primary key of identifying entity, along with extra attributes called **discriminator** to uniquely identify a weak entity
- The identifying entity set is said to **own** the weak entity set that it identifies
- The relationship associating the weak entity set with the identifying entity set is called the **identifying relationship**

- An entity set that is not a weak entity set is termed a **strong entity set**.

Weak Entity Sets-ER Diagram

- In E-R diagrams, a weak entity set is depicted via a double rectangle.
- We underline the discriminator of a weak entity set with a dashed line.
- The relationship set connecting the weak entity set to the identifying strong entity set is depicted by a double diamond
- Primary key for section – (course_id, sec_id, semester, year)

Redundant Attributes

- We model the fact that each student has an associated department using a relationship set stud_dept
- The attribute dept_name in student below replicates information present in the relationship and is therefore redundant and needs to be removed.
- BUT: when converting back to tables, in some cases the attribute gets reintroduced, as we will see later

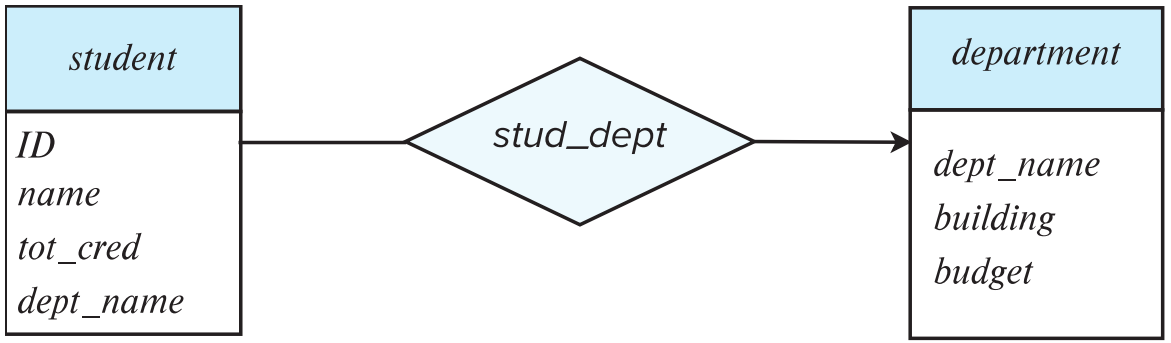

(a) Incorrect use of attribute

ER Diagram of University

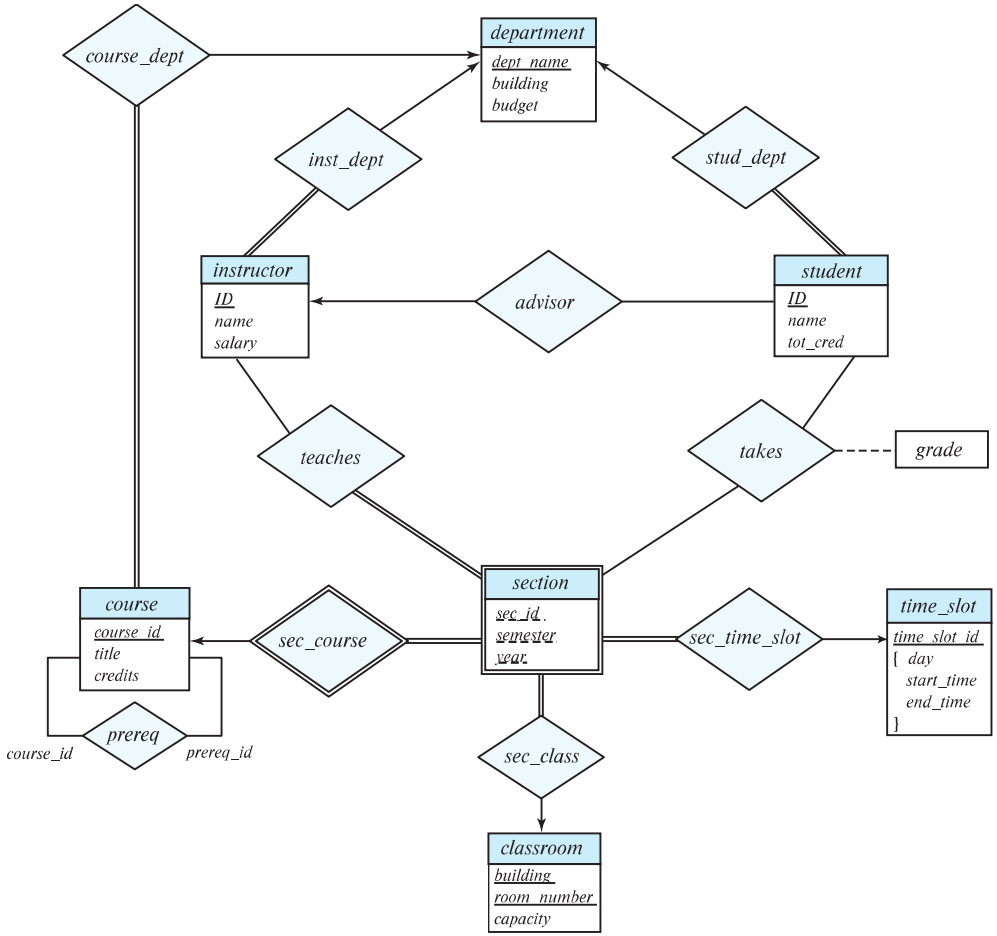

Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas* that represent the contents of the database
- A database which conforms to an E-R diagram can be represented by a collection of schemas
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set
- Each schema has a number of columns (generally corresponding to attributes), which have unique names

Representing Entity Sets

- A strong entity set reduces to a schema with the same attributes

  student(ID, name, tot_cred)

- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set

  section (course_id, sec_id, sem, year)

Composite Attributes

- Composite attributes are flattened out by creating a separate attribute for each component attribute

- Ignoring multivalued attributes, extended instructor schema is
  - instructor(ID, first_name, middle_initial, last_name, street_number, street_name, apt_number, city, state, zip_code, date_of_birth)

instructor
ID
name
first_name
middle_initial
last_name
address
street
street_number
street_name
apt_number
city
state
zip
{ phone_number }
date_of_birth
age ( )

Multivalued Attributes

- A multivalued attribute M of an entity E is represented by a separate schema EM
- Schema EM has attributes corresponding to the primary key of E and an attribute corresponding to multivalued attribute M
  - Example: Multivalued attribute phone_number of instructor is represented by a schema:
    inst_phone= (ID, phone_number)
- Each value of the multivalued attribute maps to a separate tuple of the relation on schema EM
  - For example, an instructor entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:
    (22222, 456-7890) and (22222, 123-4567)

Reduction to Relation Schemas

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set
- Example: schema for relationship set *advisor*

  advisor = (*s_id, i_id*)

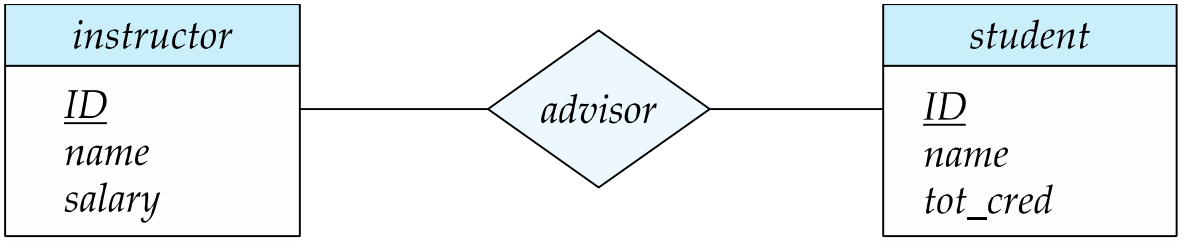

Reduction to Relation Schemas

- Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the "many" side, containing the primary key of the "one" side
- Example: Instead of creating a schema for relationship set *inst_dept*, add an attribute *dept_name* to the schema arising from entity set *instructor*

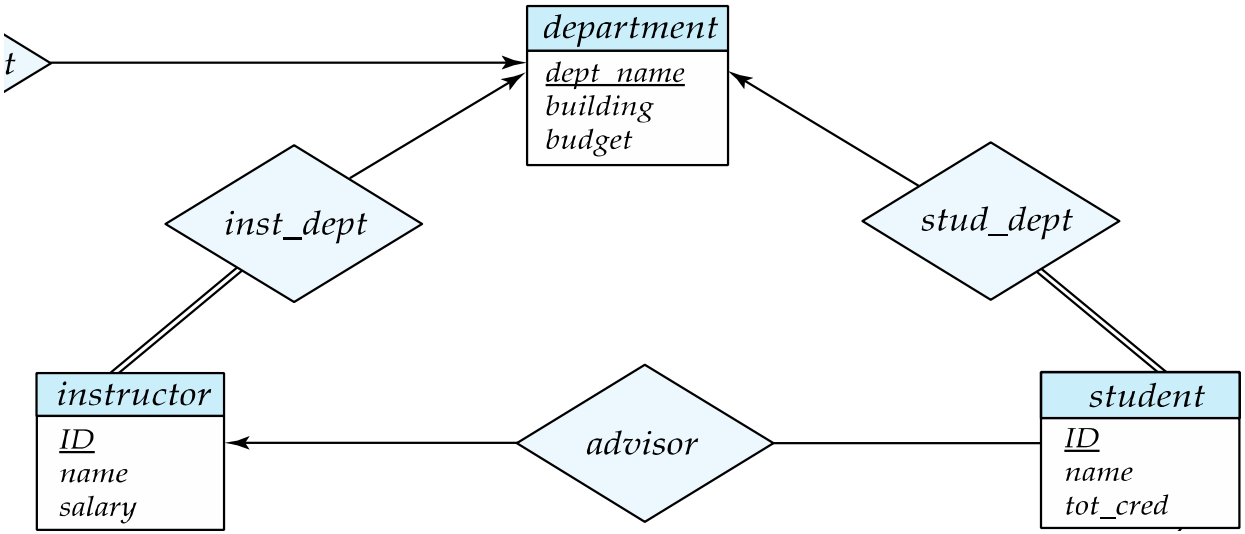

Reduction to Relation Schemas

- For one-to-one relationship sets, either side can be chosen to act as the “many” side
- That is, an extra attribute can be added to either of the tables corresponding to the two entity sets
- If participation is partial on the “many” side, replacing a schema by an extra attribute in the schema corresponding to the “many” side could result in null values

Reduction to Relation Schemas

- For one-to-one relationship sets, either side can be chosen to act The schema corresponding to a relationship set linking a weak entity set to its identifying strong entity set is redundant.
- Example: The section schema already contains the attributes that would appear in the sec_course schema

Binary Vs. Non-Binary Relationships

- Although it is possible to replace any non-binary (*n*-ary, for $n > 2$) relationship set by a number of distinct binary relationship sets, a *n*-ary relationship set shows more clearly that several entities participate in a single relationship.
- Some relationships that appear to be non-binary may be better represented using binary relationships
  - For example, a ternary relationship *parents*, relating a child to his/her father and mother, is best replaced by two binary relationships, *father* and *mother*
    - Using two binary relationships allows partial information (e.g., only mother being known)
  - But there are some relationships that are naturally non-binary
    - Example: *proj_guide*

Binary Vs. Non-Binary Relationships

- In general, any non-binary relationship can be represented using binary relationships by creating an artificial entity set
  - Replace *R* between entity sets A, B and C by an entity set *E*, and three relationship sets:
    1. R_A, relating *E* and *A*
    2. R_B, relating *E* and *B*
    3. R_C, relating *E* and *C*
  - Create an identifying attribute for *E and* add any attributes of *R* to *E*
  - For each relationship (a_i , b_i , c_i) in *R,* create
    1. a new entity e_i in the entity set *E*
    2. add (e_i , a_i) to R_A
    3. add (e_i , b_i) to R_B
    4. add (e_i , c_i) to R_C

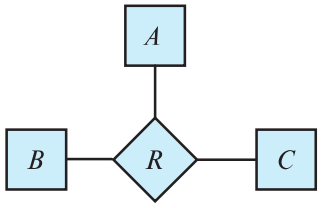

(a)

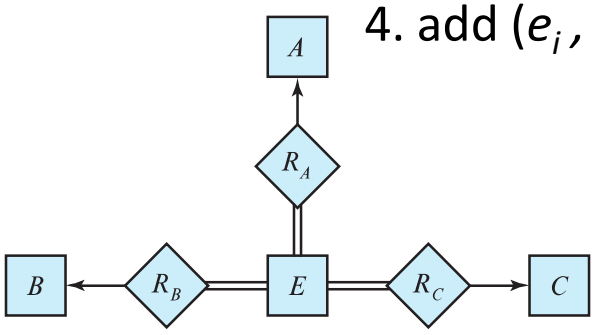

(b)

Binary Vs. Non-Binary Relationships

- Contracts(Supplier, Part, Project)
- A supplier supplies a particular part to a particular project

Supplier	Part	Project
S1	P1	J1
S1	P2	J2
S2	P1	J1

- SUPPLIER — supplies — CONTRACT
- PART — provides — CONTRACT
- PROJECT — used_for — CONTRACT

Binary Vs. Non-Binary Relationships

- CONTRACT Table

ContractID
C1
C2
C3

Binary Vs. Non-Binary Relationships

- SUPPLIER–CONTRACT relationship

Supplier	ContractID
S1	C1
S1	C2
S2	C3

Binary Vs. Non-Binary Relationships

- PART–CONTRACT relationship

Part	ContractID
P1	C1
P2	C2
P1	C3

Binary Vs. Non-Binary Relationships

- PROJECT–CONTRACT relationship

Project	ContractID
J1	C1
J2	C2
J1	C3

Acknowledgement

- Some slides and figures are collected from
  - https://www.db-book.com/